



Chapter 1: Introduction





Chapter 1: Introduction

- What Operating Systems Do
- Computer-System Organization
- Computer-System Architecture
- Operating-System Structure
- Operating-System Operations
 - Process Management
 - Memory Management
 - Storage Management
 - Protection and Security
- Computing Environments
- Open-Source Operating Systems






Objectives

- To describe the basic organization of computer systems
- To provide a grand tour of the major components of operating systems
- To give an overview of the many types of computing environments
- To explore several open-source operating systems





What is an Operating System?

- A program that acts as an intermediary between a user of a computer and the computer hardware
- Operating system goals:
 - Execute user programs and make solving user problems easier
 - Make the computer system **convenient** to use
 - Use the computer hardware in an **efficient** manner






Computer System Structure

- Computer system can be divided into four components:
 - **Hardware** – provides basic computing resources
 - ▶ CPU, memory, I/O devices
 - **Operating system**
 - ▶ Controls and coordinates use of hardware among various applications and users
 - **Application programs** – define the ways in which the system resources are used to solve the computing problems of the users
 - ▶ Word processors, compilers, web browsers, database systems, video games
 - **Users**
 - ▶ People, machines, other computers





Four Components of a Computer System



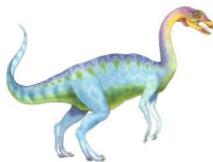


What Operating Systems Do

- Depends on the point of view
- Users want convenience, **ease of use**
 - Don't care about **resource utilization**
- But shared computer such as **mainframe** or **minicomputer** must keep all users happy
- Users of dedicated systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**
- Handheld computers are resource poor, optimized for usability and battery life
- Some computers have little or no user interface, such as embedded computers in devices and automobiles

User View





What Operating Systems Do

■ OS is a **resource allocator**

System View

- Manages all resources (CPU time, memory space, file-storage space, I/O devices,)
- Decides between conflicting requests for efficient and fair resource use

■ OS is a **control program**

- Controls execution of programs to prevent errors and improper use of the computer





Operating System Definition

- No universally accepted definition of what is part of the operating system
- “**Everything a vendor ships when you order an operating system**” is good approximation
 - But varies wildly
- “**The one program running at all times on the computer**” is the **kernel**.
 - Everything else is either a system program (ships with the operating system) or an application program.





Computer Startup


■ **Bootstrap program** is loaded when a computer is powered up or rebooted

- Typically stored in ROM or EPROM, generally known as **firmware**
- Initializes all aspects of system
 - ▶ CPU registers, device controllers, and memory contents
- Loads operating system kernel and starts execution





Computer-System Organization



- One or more CPUs, device controllers connect through common bus providing access to shared memory
- CPUs and device controllers can execute in parallel, competing for memory cycles





Computer-System Operation


- I/O devices and the CPU can execute concurrently
- Each device controller is in charge of a particular device type
- Each device controller has a local buffer
- CPU moves data from/to main memory to/from local buffers
- I/O is from the device to local buffer of controller
- Device controller informs CPU that it has finished its operation by causing an **interrupt**





Common Functions of Interrupts

- Interrupt transfers control to the appropriate interrupt service routine, through the **interrupt vector**, which contains the addresses of all the service routines
- Interrupt architecture must save the address of the interrupted instruction
- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request



An operating system is **interrupt driven**






Interrupt Handling

- The operating system preserves the state of the CPU by storing registers and the program counter
- Determines which type of interrupt has occurred:
 - **polling**
 - **vectored** interrupt system
- Separate segments of code determine what action should be taken for each type of interrupts





Interrupt Timeline





Storage Structure

- **Main memory** – only large storage media that the CPU can access directly
 - **Random access**
 - Typically **volatile**
- **Secondary storage** – extension of main memory that provides large **nonvolatile** storage capacity
 - **Magnetic disks** – rigid metal or glass platters covered with magnetic recording material
 - ▶ Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - ▶ The **disk controller** determines the logical interaction between the device and the computer
 - **Solid-state disks** – faster than magnetic disks
 - ▶ Various technologies
 - ▶ Becoming more popular






Storage Hierarchy

- Storage systems organized in hierarchy
 - Speed
 - Cost
 - Volatility
- **Caching** – copying information into faster storage system
 - Main memory can be viewed as a cache for secondary storage
- **Device Driver** for each device controller to manage I/O
 - Provides interface between controller and kernel





Storage-Device Hierarchy





Caching

- Important principle, performed at many levels in a computer (in H/W, OS, S/W)
- Information in use copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache smaller than storage being cached
 - Cache management important design problem
 - Cache size and replacement policy






I/O Structure


- After I/O starts, control returns to user program only upon I/O completion
 - Wait instruction idles the CPU until the next interrupt
 - Wait loop (contention for memory access)
 - At most one I/O request is outstanding at a time, no simultaneous I/O processing
- After I/O starts, control returns to user program without waiting for I/O completion
 - **System call** – request to the OS to allow user to wait for I/O completion
 - **Device-status table** contains entry for each I/O device indicating its type, address, and state
 - OS indexes into I/O device table to determine device status and to modify table entry to include interrupt






Two I/O Methods

Synchronous



(a)

Asynchronous




(b)





Device-Status Table







Direct Memory Access Structure

- Used for high-speed I/O devices able to transmit information at close to memory speeds
- Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention
- Only one interrupt is generated per block, rather than the one interrupt per byte



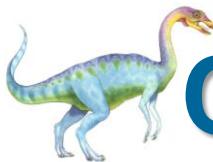


How a Modern Computer Works



A von Neumann architecture






Computer-System Architecture

- Most systems use a single general-purpose processor (PDAs through mainframes)
 - Almost all have special-purpose processors as well
- **Multiprocessors** systems growing in use and importance
 - Also known as **parallel** or **multicore** systems
 - Advantages include:
 1. Increased throughput
 2. Economy of scale
 3. Increased reliability – graceful degradation or fault tolerance
 - Two types:
 1. **Asymmetric** Multiprocessing
 2. **Symmetric** Multiprocessing






Symmetric Multiprocessing Architecture





A Dual-Core Design

- **UMA** (uniform memory access) vs. **NUMA** (non-uniform memory access) architecture variations
- Multiple chips with single cores vs. multiple **cores** on a single chip
- Systems containing all chips vs. **blade servers**
 - Chassis holding multiple independent multiprocessor systems







Clustered Systems


- Like multiprocessor systems, but multiple systems working together
 - Usually sharing storage via a **storage-area network (SAN)**
 - Provides a **high-availability** service which survives failures
 - ▶ **Asymmetric clustering** has one machine in hot-standby mode
 - ▶ **Symmetric clustering** has multiple nodes running applications, monitoring each other
 - Some clusters are for **high-performance computing (HPC)**
 - ▶ Applications must be written to use **parallelization**
 - Some have **distributed lock manager (DLM)** to avoid conflicting operations





Clustered Systems





Operating System Structure


■ Multiprogramming needed for efficiency

- Single user cannot keep CPU and I/O devices busy at all times
- Multiprogramming organizes jobs (code and data) so CPU always has one to execute
- A subset of total jobs in the system is kept in memory
- A jobs set is selected from a pool via **job scheduling**
- When a job has to wait (for I/O for example), OS switches to another job





Memory Layout for Multiprogrammed System






Operating System Structure

■ Timesharing (multitasking) is a logical extension to multiprogramming

- The CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing
- **Response time** should be short (typically < 1 second)
- Each user has at least one program executing in memory \Rightarrow **process**
- If several jobs ready to run at the same time \Rightarrow **CPU scheduling**
- If processes don't fit in memory, **swapping** moves them in and out to run
- **Virtual memory** allows execution of processes not completely in memory






Operating-System Operations

- **Interrupt driven** by hardware
- Software error or request creates **exception** or **trap**
 - Division by zero, request for operating system service
 - Other process problems include infinite loop, processes modifying each other or the operating system
- **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
 - **Mode bit** provided by hardware
 - ▶ Provides ability to distinguish when system is running user code or kernel code
 - ▶ Some instructions designated as **privileged**, only executable in kernel mode
 - ▶ System call changes mode to kernel, return from call resets it to user
- Increasingly CPUs support multi-mode operations
 - i.e. **virtual machine manager (VMM)** mode for guest **VMs**




Transition from User to Kernel Mode



■ Timer to prevent infinite loop / process hogging resources

- Set interrupt after specific period (fixed or variable)
- Operating system initializes and decrements counter
- When counter reaches zero generate an interrupt
- Set up before scheduling process to regain control or terminate program that exceeds agreed time





Process Management

- A **process** is a program in execution.
 - It is a unit of work within the system
 - Program is a **passive entity**, process is an **active entity**
- Process needs resources to accomplish its task
 - CPU, memory, I/O, files
 - Initialization data
- Process termination requires reclaim of any reusable resources






Process Management

- Single-threaded process has one **program counter** specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion
- Multi-threaded process has one program counter per thread
- Typically a system has many processes running concurrently on one or more CPUs
 - Some user processes and others system processes
 - Concurrency by multiplexing the CPUs among the processes / threads





Process Management Activities

- The operating system is responsible for the following activities in connection with process management:
 - Creating and deleting both user and system processes
 - Scheduling processes and threads on the CPUs
 - Suspending and resuming processes
 - Providing mechanisms for process synchronization
 - Providing mechanisms for process communication
 - Providing mechanisms for deadlock handling





Memory Management

- All data in memory before and after processing
- All instructions in memory in order to execute
- Memory management allows keeping several programs in memory
 - To improve CPU utilization and speed computer's response to users
- Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed





Storage Management

- OS provides uniform, logical view of information storage
 - Abstracts physical properties to logical storage unit - **file**
 - Each medium is controlled by device (i.e., disk or tape drive)
 - ▶ Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)

■ **File-System management**

- Files usually organized into directories
- Access control on most systems to determine who can access what
- OS activities include
 - ▶ Creating and deleting files and directories
 - ▶ Supporting primitives to manipulate files and directories
 - ▶ Mapping files onto secondary storage
 - ▶ Backup files onto stable (non-volatile) storage media





Mass-Storage Management

- Usually disks used to store data that does not fit in main memory or data that must be kept for a “long” period of time
- Proper management is of central importance
- Entire speed of computer operation hinges on disk subsystem and its algorithms
- OS activities
 - Free-space management
 - Storage allocation
 - Disk scheduling
- Some storage need not be fast
 - **Tertiary storage** includes optical storage, magnetic tape
 - Still must be managed – by OS or applications
 - Varies between WORM and RW





Cache Management

- Movement of information between various levels of storage hierarchy can be explicit or implicit

Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	25,000 - 50,000	5,000,000
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape


- Careful selection of cache size and replacement policy can result in greatly increased performance






Migration of Integer A from Disk to Register

- Multitasking environments must be careful to use most recent value, no matter where it is stored in the storage hierarchy



- Multiprocessor environment must provide **cache coherency** in hardware such that all CPUs have the most recent value in their cache
- Distributed environment situation even more complex (several copies of a datum can exist)





I/O Subsystem

- One purpose of OS is to hide individuality of hardware devices from the user
- I/O subsystem responsible for
 - Memory management of I/O including
 - ▶ **buffering** - storing data temporarily while it is being transferred
 - ▶ **caching** - storing parts of data in faster storage for performance
 - ▶ **spooling** - overlapping of output of one job with input of other jobs
 - General device-driver interface
 - Drivers for specific hardware devices





Protection and Security

- **Protection** – any mechanism for controlling access of processes or users to resources defined by the OS
- **Security** – defense of the system against internal and external attacks
 - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service
- Systems generally first distinguish among users, to determine who can do what
 - User identities (**user IDs**, security IDs) include name and associated number, one per user
 - User ID then associated with all files, processes of that user to determine access control
 - Group identifier (**group ID**) allows set of users to be defined and controls managed, then also associated with each process, file
 - **Privilege escalation** allows user to change to effective ID with more rights

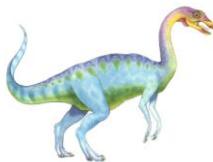




Computing Environments - Traditional

- Stand-alone general purpose machines
- But indistinct as most systems interconnect with others (i.e. the Internet)
- **Portals** provide web access to internal systems
- **Network computers (thin clients)** are like Web terminals - more security or easier maintenance
- Mobile computers interconnect via **wireless networks**
- Networking becoming ubiquitous - even home systems use **firewalls** to protect home computers from Internet attacks






Computing Environments - Mobile

- Handheld smart phones, tablet computers, etc
- What is the functional difference between them and a “traditional” laptop?
- Extra feature – more OS features (GPS, gyroscope, taking photos, and recording videos)
- Allows new types of apps like ***augmented reality***
- Use IEEE 802.11 wireless, or cellular data networks for connectivity
- Leaders are **Apple iOS** and **Google Android**





Computing Environments – Distributed

- Collection of separate, possibly heterogeneous, systems networked together
 - **Network** is a communications path, **TCP/IP** most common
 - ▶ **Local Area Network (LAN)**
 - ▶ **Wide Area Network (WAN)**
 - ▶ **Metropolitan Area Network (MAN)**
 - ▶ **Personal Area Network (PAN)**
- **Network Operating System** provides features (as file sharing) between systems across network
 - Communication scheme allows systems to exchange messages
 - Impression of a single system







Computing Environments – Client-Server

Client-Server Computing


- Dumb terminals supplanted by smart PCs
- Many systems now **servers**, responding to requests generated by **clients**
 - ▶ **Compute-server system** provides an interface to client to request services (i.e., database)
 - ▶ **File-server system** provides interface for clients to store and retrieve files





Computing Environments - Peer-to-Peer

- Another model of distributed system
- P2P does not distinguish clients and servers
 - Instead all nodes are considered peers
 - May each act as client, server or both
 - Node must join P2P network
 - ▶ Registers its service with central lookup service on network, or
 - ▶ Broadcast request for service and respond to requests for service via **discovery protocol**
 - Examples include Napster and Gnutella, **Voice over IP (VoIP)** such as Skype






Computing Environments - Virtualization

- Allows operating systems to run as applications within other OSs
 - Vast and growing industry
- **Emulation** used when source CPU type different from target CPU type (i.e. PowerPC to Intel x86)
 - Generally slowest method
 - When computer language not compiled to native code
 - **Interpretation**
- **Virtualization** – OS natively compiled for CPU, running **guest** OSs also natively compiled
 - Consider VMware running WinXP guests, each running applications, all on native WinXP **host** OS
 - **VMM** provides virtualization services






Computing Environments - Virtualization



(a)

programming
interface



(b)






Computing Environments - Virtualization

- Use cases involve laptops and desktops running multiple OSs for exploration or compatibility
 - Apple laptop running Mac OS X host, Windows as a guest
 - Developing apps for multiple OSs without having multiple systems
 - QA testing applications without having multiple systems
 - Executing and managing compute environments within data centers
- VMs can run natively, in which case they are also the host







Computing Environments – Cloud Computing


- Delivers computing, storage, even apps as a service across a network
- Logical extension of virtualization as based on virtualization
 - Amazon **EC2** has thousands of servers, millions of VMs, PBs of storage available across the Internet, pay based on usage
- Cloud compute environments composed of usual OSs plus VMMs and cloud management tools
 - Internet connectivity requires security like firewalls
 - Load balancers spread traffic across multiple applications





Computing Environments – Cloud Computing





Computing Environments – Cloud Computing

■ Many types

- **Public cloud** – available via Internet to anyone willing to pay
- **Private cloud** – run by a company for the company's own use
- **Hybrid cloud** – includes both public and private cloud components
- Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e. word processor)
- Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e a database server)
- Infrastructure as a Service (**IaaS**) – servers or storage available over Internet (i.e. storage available for backup)





Computing Environments – Real-Time Embedded Systems

- Real-time embedded systems most prevalent form of computers
 - Vary considerable, special purpose, limited purpose OS, **real-time OS**
 - Usually have little or no user interface
- Many other special computing environments as well
 - Some have OSs, some perform tasks without an OS
- Real-time OS has well-defined fixed time constraints
 - Processing **must** be done within constraint
 - Correct operation only if constraints met





Open-Source Operating Systems

- Operating systems made available in source-code format rather than just binary **closed-source**
- Counter to the **copy protection** and **Digital Rights Management (DRM)** movement
- Started by **Free Software Foundation (FSF)**, which has “copyleft” **GNU Public License (GPL)**
- Examples include **GNU/Linux** and **BSD UNIX** (including core of **Mac OS X**), and many more
- Can use VMM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms)
 - Use to run guest operating systems for exploration



End of Chapter 1

